

Manisha Subedi

Data Analyst

📍 Amadora, Lisbon, Portugal 🌐 [manisha-subedi.github.io](https://github.com/manisha-subedi)

📄 [manisubedi](#) 🔄 [manisha-subedi](#)

Experience

Data Engineer I

[Leapfrog Technology](#)

Remote
2025 – Present

- Build and maintain Python and SQL pipelines that load daily client files into a data warehouse. Add checks for duplicate rows, missing days, and delayed data.
- Use dbt to create clean tables for Power BI reports. Tests stop a failed load before it reaches a report.
- Use file hashes to keep loads idempotent, so rerunning a job does not load the same file twice.
- Work with analysts to keep each metric definition in one place, so revenue is calculated the same way in every report.

Data Engineering Intern

Bengen

Lisbon, Portugal
2025

- Wrote SQL transformations and tests for a client data warehouse. Traced data quality problems from the source files to the final report.
- Built a monitor that compares each month's row count with the usual monthly count and flags large increases.

Junior Data Analyst Intern

Company C (placeholder)

Lisbon, Portugal
2024

- Built Power BI dashboards for weekly sales and stock reports.
- Used SQL and Excel to answer ad hoc questions. Kept each metric definition with its query.

Scientific Initiation Intern

[Centro de Investigação de Montanha \(CIMO\)](#)

Bragança, Portugal
September 2023 – February 2024

- Measured physicochemical properties and analysed phenolic compounds using HPLC.
- Assessed antioxidant activity in vitro using the OxHLIA and TBARS assays.
- Processed the experimental data and prepared the final report.

Projects

recount

[Link to project](#)

- Built an MCP server for data using Python and DuckDB. When Claude requests a number, it runs the SQL, checks the tables, and returns the result with warnings.
- Checks for duplicate IDs, monthly row counts that are about twice the usual count, out-of-date data, and empty columns. Each check uses one query, and metric definitions are stored in YAML.

A warehouse on a laptop

[Link to project](#)

- Loads monthly Chicago bike-share files into DuckDB and cleans the data with dbt. The loader processed 2.3 million real rides in 25 seconds.
- Uses file hashes to identify files that have already been loaded. A dbt test fails the build when a file has twice or half the usual number of rows.

Late deliveries and bad reviews

[Link to project](#)

- Analysed 96,470 real orders from a Brazilian marketplace using SQL and DuckDB. Late deliveries made up 7% of orders and 37% of one-star reviews.
- The results pointed to promised delivery dates in some states and months rather than the sellers. In the analysis, adding three days to the promised date lowered the late-order rate from 6.8% to 4.8%.

Education

Master's degree in Food Science and Technology

NOVA School of Science and Technology (FCT NOVA)

Lisbon, Portugal
September 2025 – September 2027

Bachelor's degree in Chemical Engineering

Instituto Politécnico de Bragança

Bragança, Portugal
September 2021 – September 2024

Technical skills

SQL, Python, Databricks, DuckDB, dbt, Power BI, Tableau, Excel, pandas, Git, GitHub